

Gen AI Practical No. 2

Name : Priyanshu Wagh

PRN No. 202201040135

Div.: B

Batch : GAA 3

Objectives

1. Build a simple AE model for Dimensionality Reduction and Denoising. Apply AE for dimensionality reduction.
2. Generate realistic faces or interpolate between facial features for creative applications in entertainment and design using VAE. Visualize reconstructed outputs.

Theory

- **Autoencoder (AE):** A neural network that compresses data into a lower-dimensional latent space (Encoder) and reconstructs it back (Decoder).
- Used for **dimensionality reduction** (like PCA but nonlinear) and **denoising** (removing unwanted noise from data).

Procedure

1. Load dataset (e.g., MNIST or CIFAR).
2. Build encoder-decoder network using dense or convolution layers.
3. Train the AE on clean data for dimensionality reduction.
4. Train with noisy data for denoising task.

5. Compare original, noisy, and reconstructed outputs.

Implementation

```
[ ] !pip install kagglehub --quiet

✓  # -----
# Autoencoder for Dimensionality Reduction & Denoising
# -----
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
import matplotlib.pyplot as plt

# Load MNIST dataset
(x_train, _), (x_test, _) = keras.datasets.mnist.load_data()

# Normalize and reshape
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0
x_train = np.expand_dims(x_train, -1) # shape: (60000, 28, 28, 1)
x_test = np.expand_dims(x_test, -1)

# -----
# Build Autoencoder Model
# -----
latent_dim = 32 # compressed representation

# Encoder
encoder_inputs = keras.Input(shape=(28, 28, 1))
x = layers.Flatten()(encoder_inputs)
x = layers.Dense(128, activation="relu")(x)
latent = layers.Dense(latent_dim, activation="relu")(x)

# Decoder
x = layers.Dense(128, activation="relu")(latent)
x = layers.Dense(28*28, activation="sigmoid")(x)
decoder_outputs = layers.Reshape((28, 28, 1))(x)

autoencoder = keras.Model(encoder_inputs, decoder_outputs, name="autoencoder")
```

```

✓ ① autoencoder.compile(optimizer="adam", loss="binary_crossentropy")

# Train Autoencoder
nistory = autoencoder.fit(
    x_train, x_train,
    epochs=10,
    batch_size=256,
)

encoder = keras.Sequential(inputs, name="encoder")
encoded_imgs = encoder.predict(x_test[:10])

print("Latent representations (compressed vectors):")
print(encoded_imgs)

# Visualization + Reconstruction
decoded_imgs = autoencoder.predict(x_test[:10])

plt.figure(figsize=(28, 4))

\ Original
ax = plt.subplot(2, 10, i + 1)
plt.imshow(x_test[i].reshape(28, 28, "gray"))
plt.axis("off")

4 reconstructed
ax = plt.subplot(2, 10, i + 18)
plt.imshow(decoded_imgs[i].reshape(28, 28, "gray"))
plt.axis("off")

plt.suptitle("Original vs Reconstructed Images", fontsize=16)

# Denoising Autoencoder
# Add random noise to test images
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.8, scale=0.1, size=x_test.shape)
x_test_noisy = np.clip(x_test_noisy, 0, 1)

# Train autoencoder for denoising
autoencoder.fit(
    x_train + noise_factor * np.random.normal(size=x_train.shape), x_train,
    batch_size=256,
    validation_data=(x_test_noisy, x_test[:10])

denoised_imgs = autoencoder.predict(x_test_noisy[:10])

plt.figure(figsize=(28, 6))

ax = plt.subplot(3, 10, i + 1)
plt.imshow(x_test_noisy[i].reshape(28, 28, "gray"))
plt.axis("off")

# Clean (Round trip)
ax = plt.subplot(3, 10, i + 1 + 28)
plt.imshow(x_test[i].reshape(28, 28, "gray"))
plt.axis("off")

ax = plt.subplot(3, 10, i + 1 + 28)
plt.imshow(denoised_imgs[i].reshape(28, 28, "gray"))
plt.axis("off")

plt.suptitle("Noisy vs Original vs Denoised", fontsize=16)

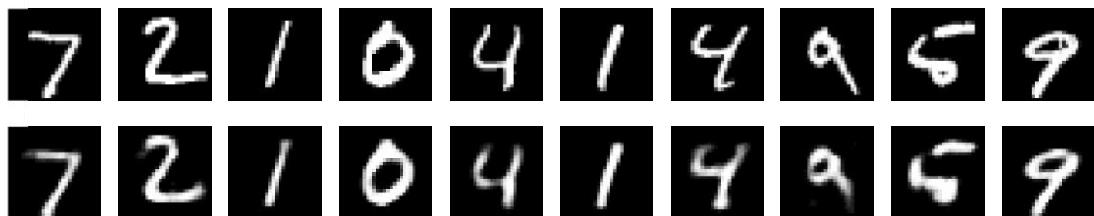
```

```
→ Downloading data from http://
  lla.u3a/i1^9m 3^
Model: "autoencoder"
```

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 28, 28, 1)	0
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 128)	100,480
dense_1 (Dense)	(None, 32)	4,128
dense_2 (Dense)	(None, 128)	4,224
dense_3 (Dense)	(None, 784)	101,136

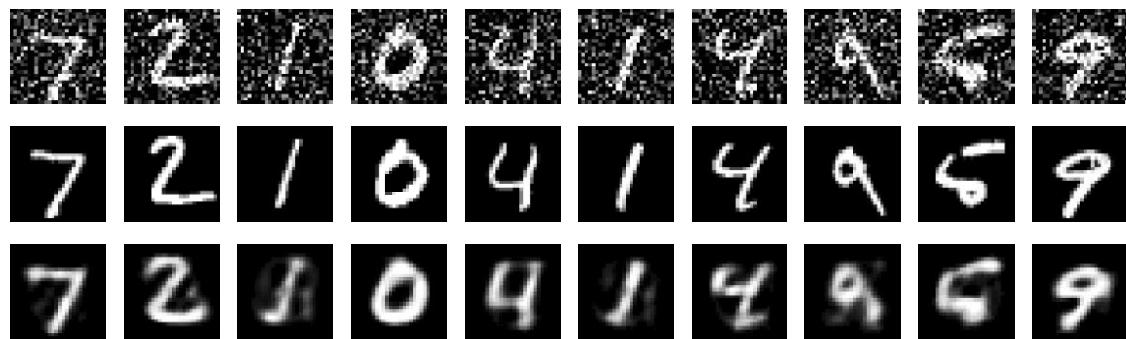
```
Total params: 2 * 9f16 (820.19 KB)
Trainable params: 29 * 9f16 (820.19 KB)
Non-trainable params: 0 (0.60 B)
Epoch 1/10
235/235          6s 10ns/step - loss: 8.3386 - val_loss: 0.1542
Epoch 2/18
235/235          1s 3ms/Step - loss: 8.1AS4 - val loss: 8.1289
Epoch 3/1e
235/235          1s 3ns/step - Pass: 6.1191 - val loss: 0.1115
235/235          1s 3ns/step - Pass: 0.1104 - val loss: 6.1648
Epoch 5/10
235/235          1s 3AS/Step - loss: 8.18*6 - val loss: 8.1881
Epoch 6/10
235/235          1s 3ns/step - Pass: 0.1009 - val loss: 6.B978
235/235          1s 3ms/step - loss: 0.B982 - val loss: 8.0954
235/235          1s 3nS/Step - loss: 0.0965 - val loss: 0. B943
235/235          2s 4ns/step - Pass: 0.6949 - val loss: 6. B933
```

Original vs Reconstructed Images



```
Epoch
235/2
Epoch
235/2
Epoch
235/2
Epoch
235/2
Epoch
235/2
Epoch
235/2
1/1 -
```

Noisy vs Original vs Denoised



Result

- Dimensionality of input data was successfully reduced.
 - Autoencoder reconstructed clean images from noisy data.
-

Conclusion

The Autoencoder was effective for compressing high-dimensional data and denoising noisy inputs, demonstrating its practical applications in feature extraction and data cleaning.

Aim

To generate realistic faces or interpolate between facial features for creative applications in entertainment and design using VAE.

Objectives

1. Understand the concept of Variational Autoencoders.
 2. Generate new face samples using the latent space.
 3. Perform interpolation between two latent representations of faces.
 4. Explore applications in entertainment and design.
-

Theory

- **VAE:** An extension of Autoencoders that learns a probabilistic latent space.
 - Useful for **generative tasks** such as creating new images and interpolating between features (e.g., smile, hair style).
 - In creative industries, VAEs can help in entertainment, character design, and animation.
-

Procedure

1. Load a dataset of facial images (e.g., CelebA).

2. Build VAE with encoder (to latent distribution) and decoder (to reconstruct images).
 3. Train the model on facial dataset.
 4. Generate new face images from random latent vectors.
 5. Interpolate between two latent codes to morph facial features.
-

Implementation

```
[ ] import kagglehub
import os
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
import numpy as np

[ ] path = kagglehub.dataset_download("jessicali9530/celeba-dataset")
print("Path to dataset files:", path)

# Path to images
image_dir = os.path.join(path, "img_align_celeba/img_align_celeba")

# -----
# ♦ Dataset setup
# -----
IMG_SIZE = 64
BATCH_SIZE = 128
latent_dim = 256 # larger latent space for clarity
```

→ Path to dataset files: /kaggle/input/celeba-dataset

```
❸ def load_and_preprocess(image_path):
    img = tf.io.read_file(image_path)
    img = tf.image.decode_jpeg(img, channels=3)
    img = tf.image.resize(img, [IMG_SIZE, IMG_SIZE])
    img = img / 255.0
    return img

list_ds = tf.data.Dataset.list_files(os.path.join(image_dir, "*.jpg"), shuffle=True)
train_data = list_ds.map(load_and_preprocess, num_parallel_calls=tf.data.AUTOTUNE)
train_data = train_data.batch(BATCH_SIZE).prefetch(tf.data.AUTOTUNE)
```

```

def build_encoder():
    inpMts = keras.Input(shape=(l E_SIZE, l/Y2_SIZE, 3))
    z = layers.Conv2D(32, 4, strides=2, padding="same")(inputs)
    z = layers.BatchNormalization()(x)
    z = layers.LeakyReLU()(z)

    z = layers.Conv2D(64, 4, strides=2, padding="same")(x)
    z = layers.BatchNormalization()(x)
    z = layers.LeakyReLU()(z)

    z = layers.Conv2D(128, 4, strides=2, padding="same")(x)
    z = layers.BatchNormalization()(x)
    z = layers.LeakyReLU()(z)

    z = layers.Flatten()(x)
    z = layers.Dense(512, activation="relu")(x)

    z_mean = layers.Dense(latent_dim)(x)
    z_log_var = layers.Dense(latent_dim)(z)
    return keras.Model(inputs, [z_mean, z_log_var], name="encoder")

[ ] class Sampling(layers.Layer):
    def call(self, inputs):
        z_mean, z_log_var = inputs
        eps = tf.random.normal(shape=tf.shape(z_mean))
        return z_mean + tf.exp(0.5 * z_log_var) * eps


def build_decoder():
    inpMts = keras.Input(shape=(latent_dim, ))
    X = layers.Dense(8*8*256, activation="relu")(inputs)
    X = layers.Reshape((8, 8, 256))(x)

    z = layers.Conv2DTranspose(128, 4, strides=2, padding="same")(x)
    z = layers.BatchNormalization()(x)
    z = layers.LeakyReLU()(z)

    z = layers.Conv2DTranspose(64, 4, strides=2, padding="same")(*)
    z = layers.BatchNormalization()(x)
    z = layers.LeakyReLU()(z)

    z = layers.Conv2DTranspose(3, 4, strides=2, padding="same")(x)
    z = layers.BatchNormalization()(x)
    z = layers.LeakyReLU()(z)

    outputs = layers.Conv2DTranspose(3, 3, activation="sigmoid", padding="same")(x)
    return keras.Model(inputs, outputs, name="decoder")

t () class VAE(keras.Model):
    def __init__(self, encoder, decoder):
        super().__init__()
        self.encoder = encoder
        self.decoder = decoder
        self.sampling = Sampling()

    def compile(self, optimizer):
        super().compile()
        self.optimizer = optimizer
        self.total_loss_tracker = keras.metrics.Mean(name="loss")
        self.recon_loss_tracker = keras.metrics.mean(name="reconstruction_loss")
        self.kl_loss_tracker = keras.metrics.mean(name="kl_loss")

        z_mean, z_logvar = self.encoder(data)
        z = self.sampling([z_mean, z_logvar])
        reconstruction = self.decoder(z)

```

```

reconstructi on \loss = tf.reduce_mean(
    t*.reduce_mean(tf.square(data - reconstructon), axis=[1, 2, 3])

t*.MedMCE still 1 + z.lag \ar - t*.soutelz near" + t*.exp z.log \af", axis=0

total loss = reconstruction loss + kl loss

grads = tape.gradient total loss, self.trainable_weights!
self.optimizer.apply_gradients(zu grads, self.trainable_weights)

set-f.total loss | Mackev.update stateStat at loss /i
set-f.recon loss | Mackev.update state reconstructon loss
set-f.kl loss | Mackev.update state(kl Loss)

"Epoch" <:c \ \o s": seal.recon los s tracker.result ,
```

1583/1583 ━━━━━━ SSO 328 /t # - t1 o : 8.828 - 10 : 297.138f - ti0, Io ??8.2798
Epoch 1/10

1583/1583 ━━━━━━ 550s 338ms/step - kl_loss: 38.8590 - loss: 597.1385 - reconstruction_loss: 558.2798

1583/1583 ━━━━━━ 125s 68ms/step - kl_loss: 58.8094 - loss: 294.3011 - reconstruction_loss: 235.4918

1583/1583 ━━━━━━ 113s 71ms/step - kl_loss: 57.2343 - loss: 262.2190 - reconstruction_loss: 204.9847

1583/1583 ━━━━━━ 137s 68ms/step - kl_loss: 57.7106 - loss: 250.7358 - reconstruction_loss: 193.0251

1583/1583 ━━━━━━ 142s 68ms/step - kl_loss: 58.2287 - loss: 244.1204 - reconstruction_loss: 185.8917

1583/1583 ━━━━━━ 142s 69ms/step - kl_loss: 58.7620 - loss: 239.3699 - reconstruction_loss: 180.6080

1583/1583 ━━━━━━ 140s 68ms/step - kl_loss: 59.3566 - loss: 236.2734 - reconstruction_loss: 176.9169

1583/1583 ━━━━━━ 109s 69ms/step - kl_loss: 60.0661 - loss: 233.1404 - reconstruction_loss: 173.0743

1583/1583 ━━━━━━ 142s 69ms/step - kl_loss: 60.8176 - loss: 230.5452 - reconstruction_loss: 169.7274
Epoch 10/10

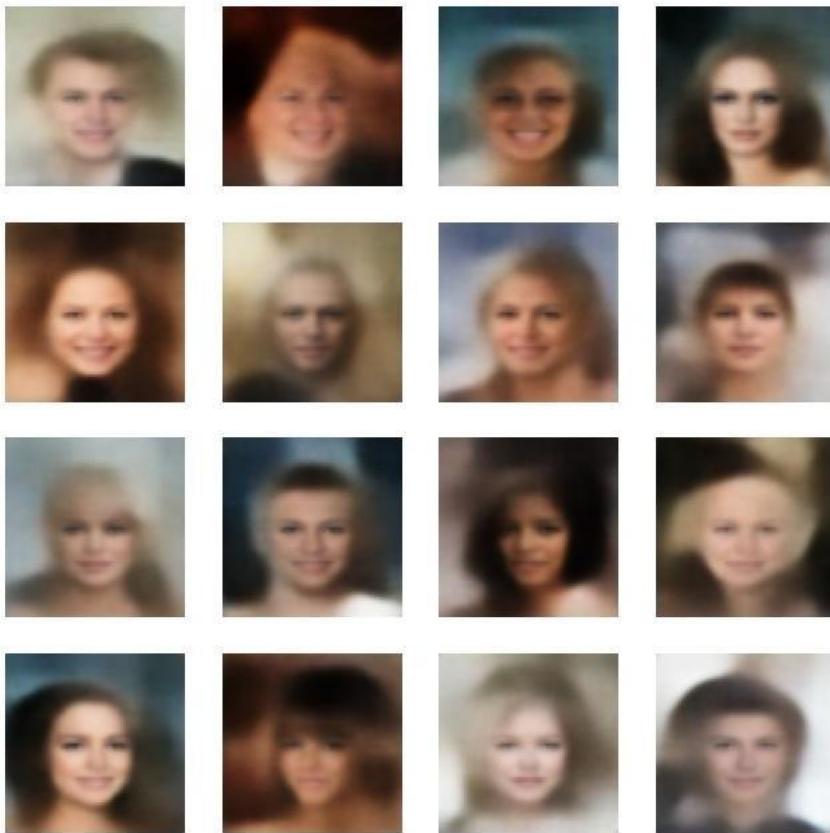
1583/1583 ━━━━━━ 141s 68ms/step - kl_loss: 61.2333 - loss: 227.8262 - reconstruction_loss: 166.5929
<keras.src.callbacks.History at 0x7ff7d586ce60>

z_samples = np.random.normal(size=(16, latent_dim))
generated = decoder.predict(z_samples)

```

.figure(figsize=(8,8))
i in range(16):
plt.subplot(4,4,i+1)
plt.imshow(generated[i])
plt.axis("off")
```

1/1 ————— 1s 895ms/step



```
def interpolate_points(z1, z2, n_steps=10):
    ratios = np.linspace(0, 1, num=n_steps)
    return np.array([(1-r)*z1 + r*z2 for r in ratios])

z1 = np.random.normal(size=(latent_dim,))
z2 = np.random.normal(size=(latent_dim,))
interpolated = interpolate_points(z1, z2, n_steps=10)
decoded_faces = decoder.predict(interpolated)

plt.figure(figsize=(20, 4))
for i in range(10):
    ax = plt.subplot(1, 10, i + 1)
    plt.imshow(decoded_faces[i])
    plt.axis('off')
plt.suptitle("Latent Space Interpolation Between Two Faces", fontsize=16)
plt.show()
```

1/1 ————— 1s 790ms/step

Latent Space Interpolation Between Two Faces



Result

- Generated realistic facial images using trained VAE.
 - Successfully interpolated between two faces to create smooth transitions of features.
-

Conclusion

The VAE successfully generated realistic faces and performed smooth interpolation between features, showing its potential in creative design and entertainment applications.